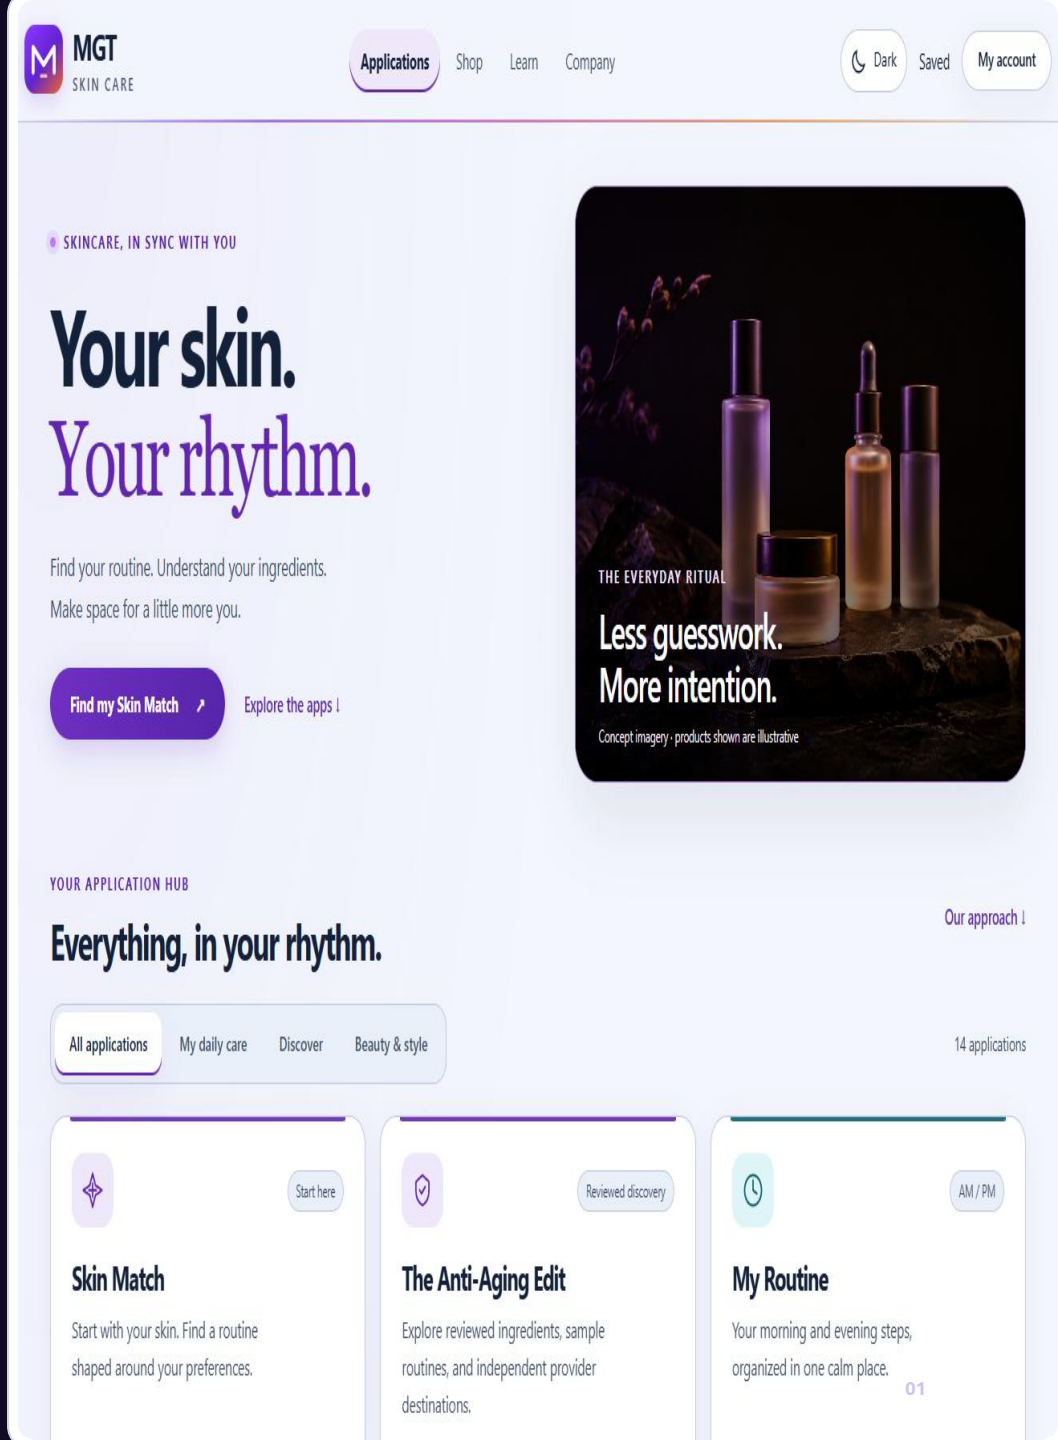
OWNER / APPROVER · GREGG CLARK

MGT Skin Care v2

Full product functions, intelligence engine, production infrastructure & launch plan

Evidence-backed portal blueprint · 3-tier commercial option · Google Cloud and Drive recovery model



EXECUTIVE READOUT

Broad, coherent product surface; production is now an evidence problem

Remaining work centers on live configuration, qualification, operations and owner approvals—not another product redesign.

14
consumer + operator surfaces

3
hosted AI roles to qualify

3
suggested subscription tiers

1
accountable owner

Built

Portal shell; profile and consent; Skin Match; routines; Anti-aging; color/style/wardrobe; Retail Edit; account/images/support/admin; Stripe contracts; local-first AI gateway; delivery scaffolding.

Must be proven

Production identity/database; webhooks; model quality/cost; approved catalog and retailers; WCAG; restore; monitoring; legal copy; domain/TLS; support and incident drills.

Decision

Approve or reject the 3-tier test. Keep live enrollment closed until prices, secrets, end-to-end payment evidence and the signed launch packet are complete.

PRODUCT SYSTEM

One consented profile drives five connected loops

Outputs combine structured profile data, reviewed knowledge, deterministic rules and bounded AI synthesis.

DISCOVER

Skin Match
Anti-aging library
Retail Edit

PLAN

AM/PM routine
Style palette
Wardrobe ideas

TRACK

Check-ins
Regimen cycles
Saved images

MANAGE

Account/privacy
Membership
Support

OPERATE

Catalog review
AI quality
Launch gates

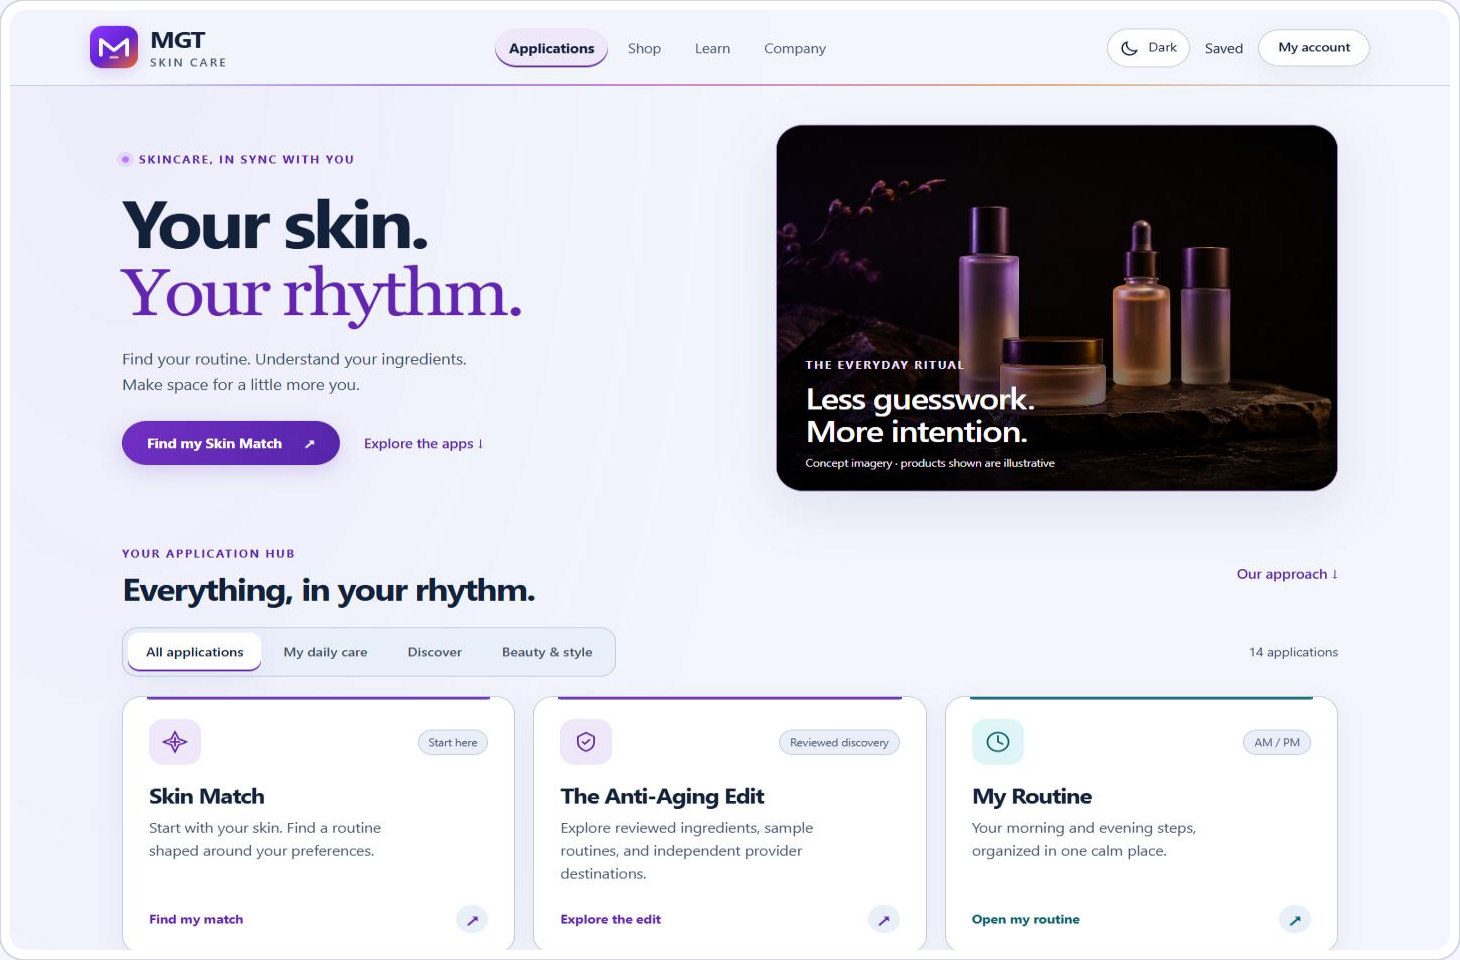
Shared profile + consent + entitlement + audit context

No AI output silently becomes a diagnosis, purchase, booking or account change.

OUTPUT SCREEN

Application hub

A progressive-disclosure front door to applications, saved state, commerce, learning and trust.



Intent first

Users begin with a goal rather than a model, provider or technical workflow.

Trust visible

Privacy, AI disclosure, support, accessibility, billing and the medical boundary are discoverable.

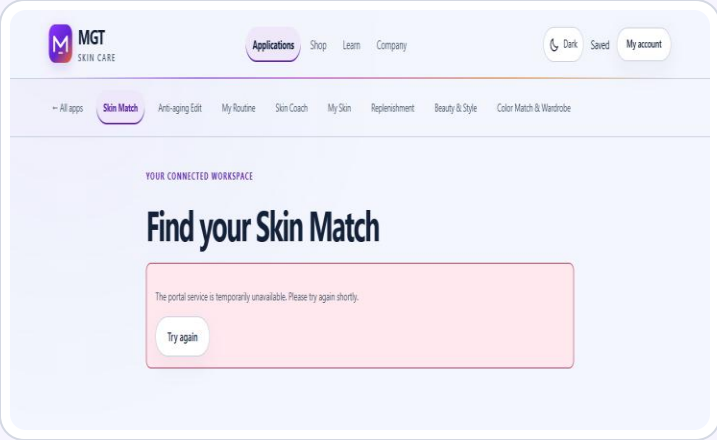
Fail closed

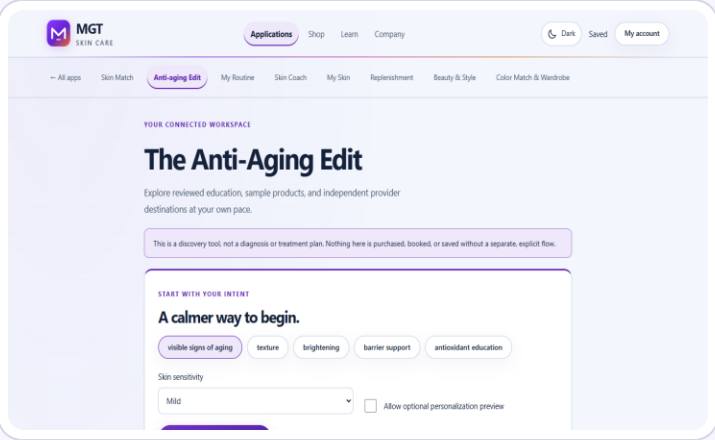
Sample data and unavailable integrations remain blocked until explicit production gates pass.

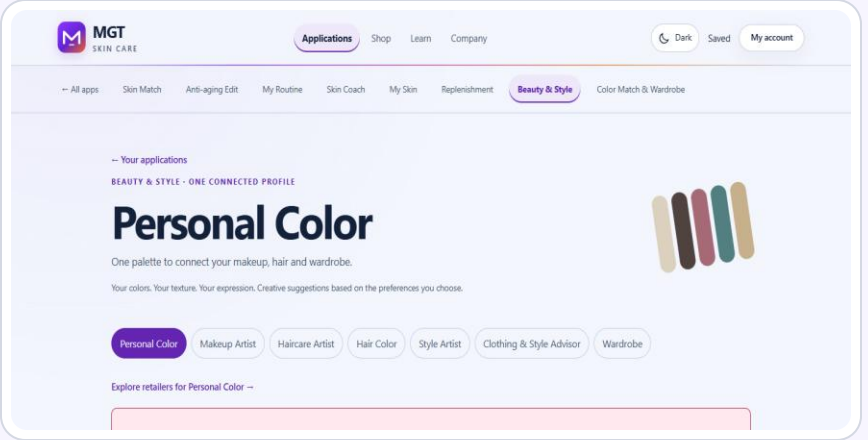
CORE JOURNEYS

Skin Match, Anti-aging, style and wardrobe

Each journey exposes structured input, explainable output and explicit save/export controls.







Skin Match → routine

Validate profile; apply avoidances and sensitivity; score approved candidates; assemble AM/PM slots; explain why; persist version and provenance.

Anti-aging boundary

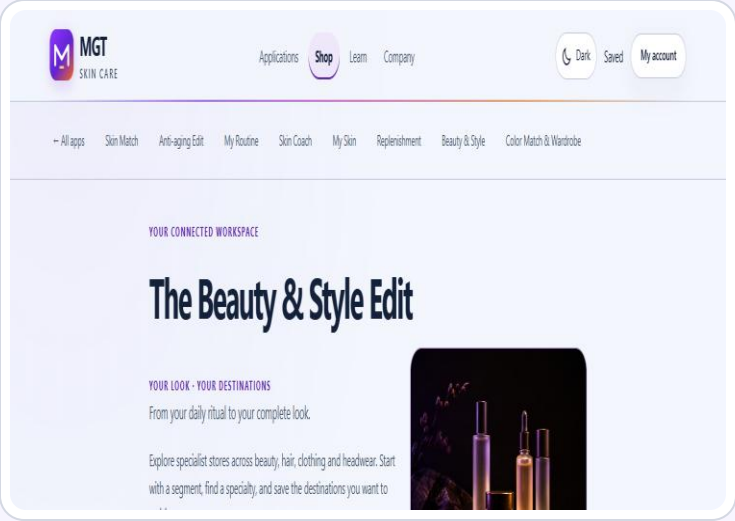
Reviewed ingredients/products/treatments/providers, routine preview, progress and reminders. Current first-load data integration remains a release gate.

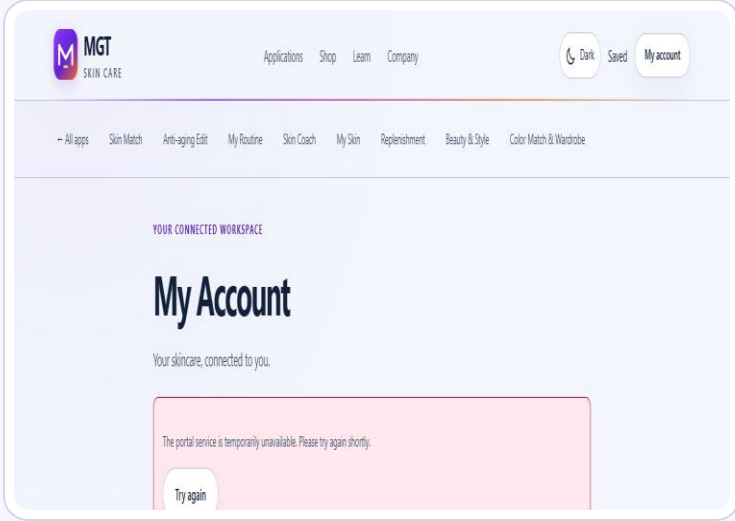
Color + wardrobe

User-confirmed appearance and comfort inputs produce a palette, explanation and exactly three practical outfit ideas—without protected-trait inference.

Retail, account, images and operator evidence

Retail stays external-commerce-only; personal data remains actor-scoped and consent-led.







Retail Edit

Allowlisted merchants, region/category filters, disclosure, saved retailers and outbound-link audit. MGT does not process product orders.

Account + images

Communication/accessibility/storage defaults; explicit upload consent; derivative images; metadata stripping; export and deletion.

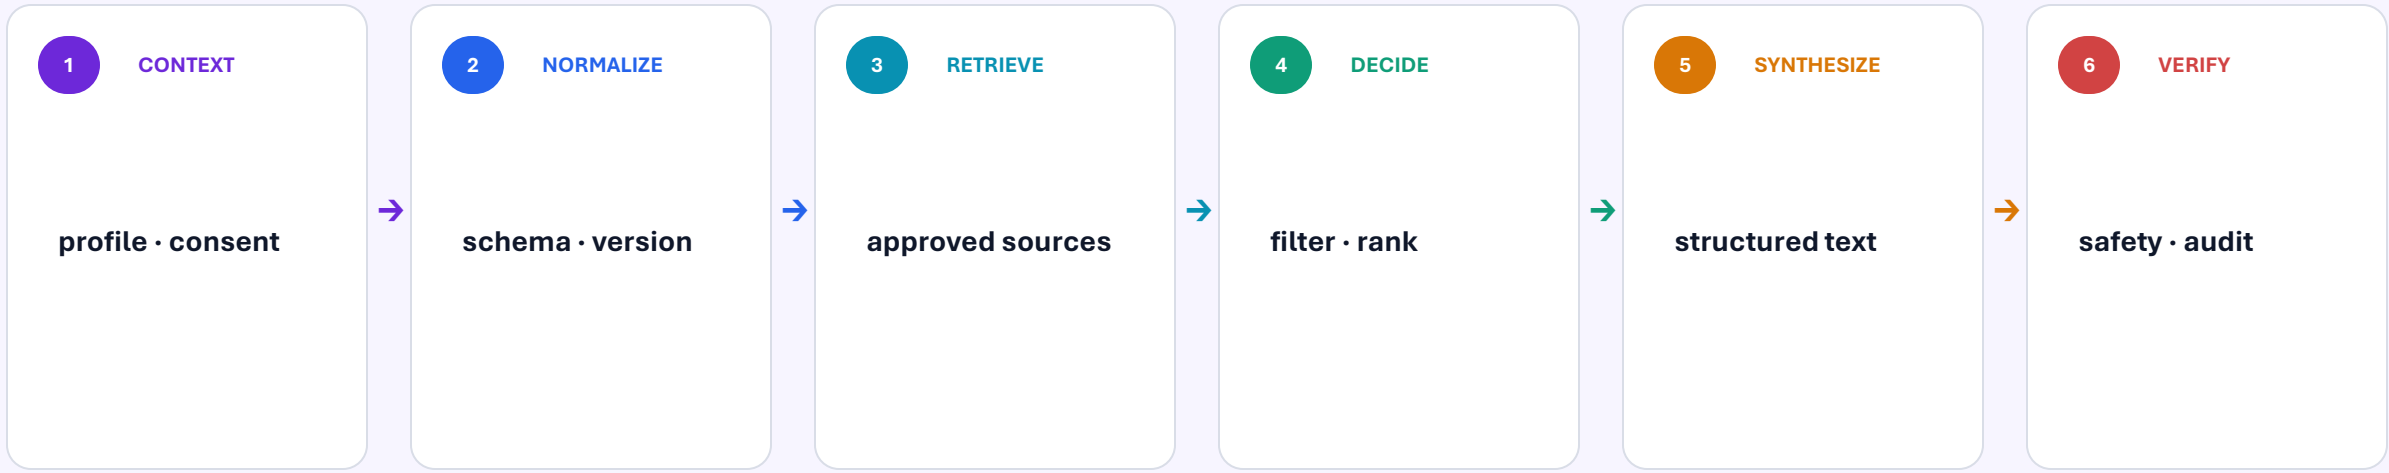
Operator console

RBAC for catalog, knowledge, AI quality/budget, readiness, support and audit—with approval separation.

ENGINE DYNAMICS

Deterministic first; AI second; evidence always

Generative text cannot override policy, scoring, catalog state or entitlement.



Integrity record

Profile version · policy/model route · prompt/template version · source IDs · confidence/uncertainty · safety flags · token/cost/latency · explicit feedback linkage.

Account, skin and care-profile inputs

Only data needed for the requested service should enter a model context.

Domain	Supported inputs	Validation / limits	Use
Account	preferred name; locale; timezone; measurement system	typed enums; change audit	language, units, reminders
Preferences	service interests; communications; accessibility; default storage	explicit and revocable	navigation, delivery and display
Skin profile	skin type; 1–3 concerns; sensitivity; age band	required structured values	matching and cautions
Routine	current routine; desired outcome; budget range	bounded choices	complexity, slots and price fit
Safety	ingredient avoidances; discomfort; optional notes	canonical IDs; length limits	hard excludes and escalation
Anti-aging	1–3 goals; sensitivity; consent; ≤6 products; ≤20 avoids	read-only boundary	preview, sources and progress
Images	purpose/category; metadata; upload consent	prepared derivative; metadata stripped; quotas	journal/reference unless separately approved

Style, retail, billing and operations

Store user facts separately from derived signals so outputs can be recomputed, explained and deleted.

Domain	Supported inputs	Validation / limits	Use
Color/style	palette, contrast, face shape, depth/undertone, hair, eyes	user-confirmed	palette and presentation ideas
Expression	occasion, texture, effort, hair, makeup, clothing	structured + short notes	bounded style brief
Access needs	coverage, headwear, climate, fastening, comfort, sensory	preference-led; revocable	filter impractical outputs
Wardrobe	item ID/name/category/color/notes; image reference	≤50 actor-scoped items	three-item combinations and gaps
Retail	region, category, merchant, saved retailer, outbound click	allowlist; no in-MGT order	discovery and disclosure audit
Billing	plan/cycle, Stripe IDs, status, period end	webhook authoritative; idempotent	entitlement and self-service
Operations	catalog approval/source/reviewer; route; latency/cost; feedback	role-separated audit	release gates and QA

OUTPUT CONTRACT

Maximize utility without sacrificing integrity

The best output is traceable, current, personalized and safely actionable—not merely long.

Profile vector

Versioned normalized features, missingness, user-vs-derived lineage.

Match result

Rank, score components, exclusions, alternatives, budget and availability.

Routine / look

Ordered steps or three ideas; timing, comfort and effort constraints.

Explanation

Rationale, source IDs, confidence, uncertainty and next question.

Safety packet

Cautions, out-of-scope detection and human escalation.

Evidence record

Model/prompt/catalog versions, cost, latency, checks and feedback.

Ask for missing high-impact fields → retrieve approved records → generate JSON-schema output → validate → show provenance/uncertainty → learn from explicit feedback.

BACKEND INFRASTRUCTURE

Current stack and production destination

Small logical services, isolated identities and fail-closed external dependencies.

Experience	Next.js 15 portal · Expo/Android contract · accessible design system
Edge & identity	Caddy / managed HTTPS · Supabase Auth + Google OAuth · CSRF · rate limits · RBAC
Application	Express + TypeScript API · domain/shared packages · worker · typed validators
Data	SQLite local/demo → PostgreSQL/Supabase or Cloud SQL prod · migrations · audit
Intelligence	AI gateway · deterministic match/style engines · retrieval · routing · eval telemetry
Commerce/integrations	Stripe Checkout/Portal/Webhooks · Drive test/evidence · retailer allowlist
Delivery/operations	Docker/Compose · GitHub Actions · immutable images · monitoring · backup/restore

Current local-first registry; live qualification pending

The repository defines routes and fallbacks, but available credentials/model tags and representative evals are still required.

Role	Configured model	Workload	Strength	Status
Local volume	llama3.2:3b	classification, rewrite, FAQ	private; low marginal cost	model tag required
Local reasoning	deepseek-r1:8b	reasoning fallback	local control	model tag required
Local vision	llava:latest	consented image description	no hosted transfer	not skin diagnosis
Embeddings	nomic-embed-text	retrieval/vector search	local and inexpensive	index evidence required
Hosted synthesis	gpt-5.6-sol	premium consultation	structured synthesis	credential/live eval pending
Qualification	GPT-6 Luna / 6.1 Sol / Astra	volume / balanced / exceptional	capability-cost ladder	registered; no switch approved

RECOMMENDED AI ALTERNATIVE

Use a measured router—not one universal model

Start hosted for operational simplicity; retain deterministic engines and add local inference only when scale/privacy economics justify it.

Route	Suggested model	Examples	Effort	Economics / guardrail
Fast path	GPT-6 Luna	classification, normalization, short variants	none / low	$0.10 in / $0.50 out per MTok; schema + audit
Core synthesis	GPT-6.1 Sol	profile, routine/style explanation, support	low / medium	$2 in / $10 out; default candidate after eval
Escalation	GPT-6 Astra	complex conflicts and operator investigation	medium / high	$10 in / $50 out; budget cap + review
Retrieval	local embedding + approved records	catalog and knowledge grounding	n/a	versioned index; no open-web medical synthesis
Vision	separate consented route	quality/non-diagnostic description	task-specific	redaction; short retention; no trait inference

Promote only when representative evals pass: task success · safety · p50/p95 latency · cost per successful task · refusal quality · schema validity.

PAYMENTS & ENTITLEMENT

Stripe bills; MGT projects access

Never grant access from the redirect. Provision only after verified, idempotently processed server-side events.

1

Choose

plan + cycle

→

2

Checkout

hosted session + terms

→

3

Webhook

signature + event ID

→

4

Entitle

subscription → access

→

5

Portal

payment, invoice, cancel

Configured

Checkout Session, monthly/annual mapping, trial support, signed webhooks, billing activity, entitlement endpoint, retry and Customer Portal.

3-tier change

Add Guided and Studio monthly/annual price families, update entitlement matrix and define migration/downgrade behavior. Free is non-billable.

Manual proof

Live IDs in Secret Manager; webhook replay; test trial, payment failure, cancel, refund/dispute and plan change; retain evidence.

Recovery first; Drive visibility second

A complete backup is a recoverable data/object system plus a readable evidence archive—not a single Drive folder.

Cloud runtime

1. **Artifact Registry: immutable images**
2. Cloud Run: web + API
3. Job/worker: scheduled tasks
4. Secret Manager + service IDs
5. Logging, Monitoring and alerts

Recovery plane

1. **Choose Supabase or Cloud SQL primary**
2. PITR + automated snapshots
3. Versioned GCS object/export bucket
4. KMS + retention/immutability
5. Quarterly isolated restore drill

Drive evidence archive

1. **Launch packets and approvals**
2. Redacted inventory + checksums
3. Restore reports; no raw prod DB
4. PDF/PPTX and runbooks
5. Gregg-owned access review

MARKET ANCHORS

Price MGT as digital guidance—not prescription care or shipped product

Adjacent public offers are directional anchors, not direct like-for-like software comparables.

Offering	Public price	Cadence	Included value	MGT implication
Curology Custom Formula Rx	$29.95 / month	subscription	Rx formula + provider personalization	stay below Rx care unless clinical services added
Musely Anti-Aging Cream Rx	$35 / month	2-month supply	Rx product + doctor support	ceiling for non-clinical guidance
PROVEN Skincare System	$142.77 / shipment	45–60 days	personalized physical system	high value includes products
MGT proposed	$0 / $14.99 / $29.99	monthly + annual	software guidance, tracking, style, discovery	test willingness-to-pay and retention

Pricing requires owner approval, customer interviews and controlled conversion/retention testing.

SUGGESTED 3-TIER MODEL

Trust at Free; continuity at Guided; breadth and depth at Studio

Commercial alternative to the repository's locked two-tier release contract—requires owner sign-off and code/config changes.

FREE

$0

Try

Skin profile; sample match; basic color; limited saves; retail discovery; privacy/export.

GUIDED

$14.99/mo

$149.99/yr

Full routine, anti-aging education, progress, wardrobe, images, history, standard AI and reminders.

~2 months free annually

STUDIO

$29.99/mo

$299.99/yr

Higher AI allowance, cross-domain synthesis, advanced exports, priority support, early features and guest access.

~2 months free annually

Test price × annual anchor × AI allowance × trial. Guardrails: support load, AI cost/member, refunds and safety incidents.

Eight workstreams, each closed by reproducible evidence

No workstream is complete merely because a setting exists or a test passed once.

1. Product

scope, copy, catalog, retailers, benefits

Evidence: acceptance + screens

2. Identity/privacy

prod tenant, OAuth, roles, consent, retention

Evidence: access + privacy test

3. Data/backup

DB, migrations, RLS, PITR, GCS, Drive evidence

Evidence: restore + checksum

4. AI/quality

routes, prompts, evals, budget, red-team

Evidence: eval report + approval

5. Payments

products/prices, Checkout, Portal, webhooks, tax

Evidence: Stripe lifecycle proof

6. Reliability

Cloud Run, scaling, SLOs, monitoring, on-call

Evidence: load + alert drill

7. Access/support

WCAG, keyboard, screen reader, support mailbox

Evidence: audit + ticket drill

8. Release

CI/CD, digest, rollback, change log

Evidence: staging sign-off

MANUAL OWNER INPUTS

Gregg Clark is owner, approver and launch operator

Legal/business facts, secrets and external-console settings cannot be inferred from the repository.

System	Manual input	Evidence	Go / no-go
Business/legal	entity, address, support, terms/privacy, AI/medical disclosure	approved docs + dates	links live
Google Cloud	project, billing, region, domain, IAM, quotas	inventory + access review	least privilege + alerts
Data/storage	provider, retention, RPO/RTO, restore approver	backup + restore result	objective met
Stripe	verification, bank/tax, products/prices, portal, webhook	test + live receipts	paid state matches access
AI/models	API project/key, spend, allowed data, availability	golden eval + cost report	targets pass
Google Drive	folder, owner/group access, retention, folder IDs	upload + readback	no DB dump/secrets
Launch	price, benefits, trial, support, date, rollback authority	signed packet	Gregg approval recorded

LAUNCH SEQUENCE

Realistic six-week evidence path after decisions and accounts are available

Schedule is a planning proxy, not a promise; external verification remains on the critical path.

W1

DECIDE

pricing, provider, data boundary, RPO/RTO, legal

W2

CONFIGURE

identity, database, secrets, Stripe, Google projects

W3

INTEGRATE

migrations, webhooks, Cloud Run, GCS/Drive

W4

QUALIFY

AI, accessibility, security, load and failure

W5

REHEARSE

restore, rollback, incident, support and billing

W6

RELEASE

acceptance, signed packet, canary, observe

Compute allocation

Product/legal 20% · cloud/data 25% · payments/identity 15% · AI qualification 20% · QA/accessibility/security 15% · operations 5%.

Stop conditions

No restore evidence; unverified webhook; critical access/security defect; missing legal links; unapproved live catalog; AI regression; unsupported rollback.

Move from build to controlled production qualification

Four explicit approvals start the final path

01 **Commercial model**

Keep two tiers or implement Free / Guided / Studio.

02 **Data path**

Supabase + Cloud Run, or Cloud SQL + Cloud Run—not both primary.

03 **AI budget**

Authorize Luna, Sol and Astra evaluation and data boundary.

04 **Launch window**

Name date, support, reviewers and rollback authority.